

Final Readiness Assessment

Paper: “Institutional Observability Under Scaled AI Governance” **Date:** June 4, 2026

The real question

Can an editor send this to review?

Yes.

Would I predict acceptance?

Not yet. The paper is theoretically sound and the argument is clear, but it has vulnerabilities that a determined reviewer will find. Here they are, honestly.

What’s solid

- **The core idea is genuinely novel.** Oversight saturation as a structural consequence of throughput asymmetry is not something I’ve seen formalized this way before. The SEDI index is clever — computable from public records, no proprietary access needed. That’s a real contribution.
 - **The math checks out.** Kingman’s heavy-traffic approximation giving 2.71x is correct for the stated parameters. The degradation function $D(\rho)$ is well-behaved. The state-space model is internally consistent.
 - **The figures exist and are clean.** All 5 figures are generated from actual running code. No mockups. No placeholders. The simulation produces real outputs.
 - **The policy implications are concrete.** Three governance-by-design principles that name specific actors and specific actions. Not vague recommendations.
 - **The GitHub repo is live.** <https://github.com/SilverCoin256/oversight-saturation-sedi> — simulation code, figures, manuscript PDF, cover letter. A reviewer can clone it and run `python simulation/figure_generation.py` to reproduce everything.
-

What’s vulnerable

1. No empirical validation

This is the biggest weakness and there’s no way around it. The paper is theoretical + simulation. The calibration uses published reports (Vanderbilt, PHTI), not primary data. A reviewer who demands empirical evidence will not be satisfied by the limitations section alone.

Mitigation: The limitations section is honest about this. The paper is positioned as theoretical groundwork. But if the reviewer is an empiricist, this will be a problem.

2. The Monte Carlo claims are specific

The paper states: 120 runs, 95% CIs, $R^2 > 0.95$, specific SEDI values at specific deployment stages. These claims are backed by actual running code in the repo. If a reviewer asks for the code, you can point them to it and they can run it themselves.

But: if the code didn't exist, this would be the single biggest risk in the paper. A reviewer asking "please provide simulation code" and getting nothing back is how papers get rejected with prejudice. The code exists. Make sure it stays accessible.

3. Parameter calibration is thin

$k=3.50$ and $\rho_{oc}=0.60$ come from one 1997 paper (Parasuraman & Riley). That's the canonical reference for alert fatigue, but it's still one source. A reviewer can reasonably ask: why these values? What happens if $k=2.0$ or $k=5.0$?

Mitigation: The limitations section acknowledges this. A sensitivity analysis appendix would help but doesn't exist yet.

4. Three AAAI 2026 workshop papers in the references

iCRAFT, P2T, and RG-PSO are all from the same workshop. They're cited substantively (not just name-dropped), but a reviewer might question whether workshop papers belong in a journal submission's reference list.

Mitigation: The core references (Meyer & Rowan, Kingman, Parasuraman & Riley, Bovens, Raji et al.) are established. The workshop papers provide architectural context for deployed systems. It's defensible but not ideal.

5. The writing style is a hybrid

Parts of the paper read like a CS workshop paper (dense equations, $O(N)$ notation). Parts read like a policy brief. The introduction and conclusion now have personal voice, but the middle sections are still technical. A social science reviewer might find the math density off-putting.

Mitigation: Explanatory prose surrounds the equations. But the paper is what it is — a technical governance paper.

Figure verification

All 5 figures confirmed present and generated from running code:

Figure	File	Size	Status
1: Saturation Curve	fig_saturation.pdf	35 KB	[PASS]
2: Divergence Plot	fig_divergence.pdf	31 KB	[PASS]
3: Routing Architecture	fig_routing.pdf	27 KB	[PASS]
4: Review Depth KDE	fig_depth_kde.pdf	74 KB	[PASS]
5: SEDI Degradation	fig_sedi.pdf	33 KB	[PASS]

They exist in both the submission `figures/` directory and the GitHub repo.

GitHub repo status

URL: <https://github.com/SilverCoin256/oversight-saturation-sedi>

Structure matches the recommended layout:

<code>simulation/</code>	<code><- all Python code (6 files)</code>
<code>figures/</code>	<code><- all 5 figure PDFs</code>
<code>manuscript/</code>	<code><- paper PDF + cover letter PDF</code>
<code>README.md</code>	<code><- human-written, no badges</code>
<code>requirements.txt</code>	<code><- numpy, scipy, pandas, matplotlib</code>
<code>LICENSE</code>	<code><- MIT</code>
<code>rules.md</code>	<code><- project rules</code>
<code>prof_rules.md</code>	<code><- mentor advice</code>

The repo is public. A reviewer can clone it and reproduce everything in under 5 minutes.

What to do before submitting

1. **Read the paper out loud once.** Catch any remaining LLM-sounding sentences. The introduction and conclusion are good now. Check the methods sections.
 2. **Verify the .docx opens correctly.** The LaTeX-to-docx conversion can mangle figure placement. Open it in Word and check.
 3. **Make sure the GitHub repo stays up.** If the repo goes private or gets deleted, the simulation claims become unverifiable.
 4. **Consider adding a sensitivity analysis.** Even a simple table showing SEDI values for k in $\{2, 3.5, 5\}$ and ρ_{hoc} in $\{0.5, 0.6, 0.7\}$ would preempt the parameter calibration criticism.
-

Bottom line

This paper can go to review. It has a real contribution, working code, clean figures, and honest limitations. It will get criticisms — every paper does — but the core argument is defensible.

The editor's decision to send it to review depends on whether they find the topic important enough to overlook the lack of empirical validation. Given the current interest in AI governance and the novelty of the saturation framing, I think they will.

Acceptance depends on who reviews it. A theoretically-inclined reviewer will engage with the model. An empiricist will ask for data you don't have. That's the gamble.