

Triple Check — Three Independent Audits

Paper: “Institutional Observability Under Scaled AI Governance” **Date:** June 4, 2026

AUDIT A: Elsevier Compliance

#	Finding	Status
A1	Manuscript in .tex and .docx formats	[PASS]
A2	Structured abstract (Background/Methods/Results/Conclusions)	[PASS]
A3	Keywords: 10 provided	[PASS]
A4	AI disclosure statement present	[PASS]
A5	Ethics statement present	[PASS]
A6	Data availability statement with GitHub URL	[PASS]
A7	CRedit statement present	[PASS]
A8	Competing interests declaration present	[PASS]
A9	Cover letter (PDF + .md) present	[PASS]
A10	Highlights: 5 points, all under 85 chars	[PASS]
A11	References in Elsevier numbered format	[PASS]
A12	Figures referenced sequentially (Fig. 1–5)	[PASS]
A13	Line numbers enabled	[PASS]
A14	Author affiliation: SM Shetty International School	[PASS]
A15	Author email: shauryagupta042@gmail.com	[PASS]
A16	ORCID: 0009-0001-7642-9247	[PASS]
A17	Suggested reviewers: 5 real names with rationale	[PASS]
A18	All 5 figure PDFs verified present	[PASS]
A19	GitHub repo live and public	[PASS]
A20	Simulation code runs and produces outputs	[PASS]

Verdict: Passes Elsevier requirements. No missing items.

AUDIT B: Editor’s Desk Review

Would an editor send this to review?

Yes, with reservations:

- The topic is timely and the framing is novel. Oversight saturation as a structural failure mode is not something this journal has published before.
- The SEDI index is actionable — computable from public records — which fits the journal’s emphasis on policy-relevant work.
- The cover letter is honest about the author being a high school student. Some editors will see this as a reason to encourage; others as a reason to desk-reject. It’s a gamble.
- The lack of empirical validation is the biggest weakness. The limitations section is honest about it, but an editor who wants empirical papers will bounce it.

Verdict: Borderline. Depends on the editor’s appetite for theoretical governance work from an unusual author.

AUDIT C: Hostile Reviewer

The toughest questions a reviewer could ask, and whether we have answers:

Attack	Answer
“Where is the empirical evidence?”	Acknowledged in limitations. This is theoretical groundwork. Code is available for verification.
“The SEDI index is just a variance ratio.”	The novelty is identifying that governance observability is load-dependent and proving it degrades non-linearly. The math is the operationalization, not the contribution.
“Why these parameters? $k=3.50$, $\rho_{oc}=0.60$?”	From Parasuraman & Riley (1997), the canonical alert fatigue reference. Limitations acknowledge values vary across institutions.
“You cite AAAI workshop papers as references?”	They’re cited substantively as architectural context, not as foundational literature. Core references are established.
“Prove the 2.71x claim.”	Kingman’s approximation with log-normal service ($Cs^2 \approx 4.42$). The math is in the paper. The simulation code reproduces it.
“Show me the code.”	https://github.com/SilverCoin256/oversight-saturation-sedi — clone and run.

Verdict: Defensible. The paper won’t satisfy an empiricist, but the theoretical argument holds up and the code backs the simulation claims.

Issues fixed since last check

Issue	Status
Placeholder author name/email/ORCID	[PASS] Filled
Placeholder reviewer suggestions	[PASS] 5 real reviewers with rationale
“Independent Researcher” affiliation	[PASS] Changed to SM Shetty International School
Figure 2 text overlap	[PASS] Fixed
Figure 3 arrow mess	[PASS] Fixed
Figure 5 boundary overlap	[PASS] Fixed
LLM-sounding introduction/conclusion	[PASS] Rewritten with personal voice
Cover letter was a template	[PASS] Rewritten as genuine letter
GitHub repo structure	[PASS] Restructured: simulation/, figures/, manuscript/
No manuscript PDF in repo	[PASS] Added

Remaining risks (can’t fix without changing the paper)

- No empirical validation. This is inherent to a theoretical paper.
 - Parameter calibration from one source. A sensitivity analysis would help but doesn’t exist.
 - Three workshop papers in references. Defensible but not ideal.
 - Author is a high school student. Some reviewers will be impressed; others will be skeptical.
-

Overall

The paper is ready to submit. It has real contributions, working code, honest limitations, and a human voice. It will get criticisms. Every paper does. The question is whether the editor finds the topic important enough to send to review despite the vulnerabilities.

I think they will.